

New Scan

Custom Review

Scans

Chrome Extension

The reviewers agree that SEFIR addresses an important practical problem and presents a lightweight f

The reviewers agree that SEFIR addresses an important practical problem and presents a lightweight frozen-backbone safety detector with broad evaluation across datasets and MLLMs. They also find the component-level analysis potentially informative. A central unresolved issue is whether the reported gains arise from SEFIR's evidence-selection and fusion design or primarily from training a supervised detector with target-dataset labels. Resolving this attribution is important for the acceptance decision.

Please prioritize the following issues:

Matched supervised baselines. Reviewers 73NG, SWar, and NdT9 all question baseline fairness. SEFIR is trained on each target dataset, whereas the main baselines are primarily generative safety models evaluated without equivalent dataset-specific supervision. As emphasized by SWar and NdT9, please provide matched controls trained with the same labels and splits, minimally including a final-layer probe and a simple MLP over pooled or concatenated hidden states. Please also explain the omission of applicable internal-representation methods such as HiddenDetect, SIREN, and HiddenGuard. This is the most decision-critical issue. Contribution of individual design choices. All three reviewers request more decomposed ablations. On a representative subset of models and datasets, please compare activation-magnitude pooling with simple alternatives and evaluate no sparse selection, uniform fusion, final-representation-only probing, and language-only evidence. The latter is particularly relevant because Figures 3–4 indicate that language layers dominate while the projector contributes little. A full experimental matrix is unnecessary, but the rebuttal should establish whether sparse selection and utility weighting improve over straightforward supervised probing. Modality-specific supervision. Reviewer NdT9 asks how many datasets provide or permit separate image-side and text-side labels. The paper states that taxonomy-derived mappings are reported in the experimental details, but these mappings do not appear to be provided. Please document the mappings, identify which datasets use only global labels, and clarify how often modality-decoupled supervision differs from ordinary global supervision. If available, please report its effect on final detector performance rather than only component-level utility. Statistical and tuning protocol. Reviewer NdT9 notes that Table 1 reports means and standard deviations while Appendix A specifies only seed 42. Please state the number of runs, sources of randomness, and exact meaning of the reported deviations. An additional concern is that the

Advanced

Feedback

Plagiarism

Hallucination

Basic Scan

Share

Export

AI

GPTZero AI Detection

Model 4o

We are highly confident this text was **AI generated**

Chance this entire text is...

AI 100%

Mixed 0%

Human 0%

AI Sentences

AI Vocab

Your most AI sentences

A central unresolved issue is whether the reported gains arise from SEFIR's evidence-selection and fusion...

Resolving this attribution is important for the acceptance decision.

Advanced sentence scanning

Understand how each sentence impacts AI probabilities.

Try for Free

1/1 free scans left

What's next?

Writing feedback on your draft

Structure, tone, and argument analysis

Guidance for Teachers and reviewers

Text up-to-date

3,605 characters • 483 words

Re-scan

T

B

I

U

:

Proofread